Question 1:

How long does it take to load 200B bf16 model into the systolic array of 32 TPU v4p from HBM

TPU v4p HBM to Tensorcore bandwidth

$= 1.2 \times 10^{12}$

Model size $= 2 \times 10^{9}$

$\text{Sharded\_model} = \frac{2 \times 2 \times 10^{9}}{32} \times 100$

$$\frac{\text{Sharded\_model}}{\text{HBM to Tensorcore BW}} = \frac{\frac{2 \times 2 \times 10^{9} \times 100}{32 \times 16}}{1.2 \times 10^{12}} = \frac{2 \times 1 \times 100}{1.2 \times 16 \times 10^{3}} \text{ sec}$$

$$= \frac{2 \times 1 \times 100}{19.2 \times 10^{3}}$$

$$= 0.052 \times 10^{-3} \times 100$$

$$= 2 \times 5.2 \times 10^{-5} \text{ sec}$$

$$= 10.4 \times 10^{-5} \times 100$$

$$= 10.4 \times 10^{-3}$$

$$= 10 \text{ ms}$$

Question 2: TPU v5e pod

① How many CPU hosts are there?

TPU v5e pod size = 16 x 16

Every 4X2 contains a CPU

So = $\frac{16}{4} \times \frac{16}{2} = 4 \times 8 = 32$ CPU hosts ✓

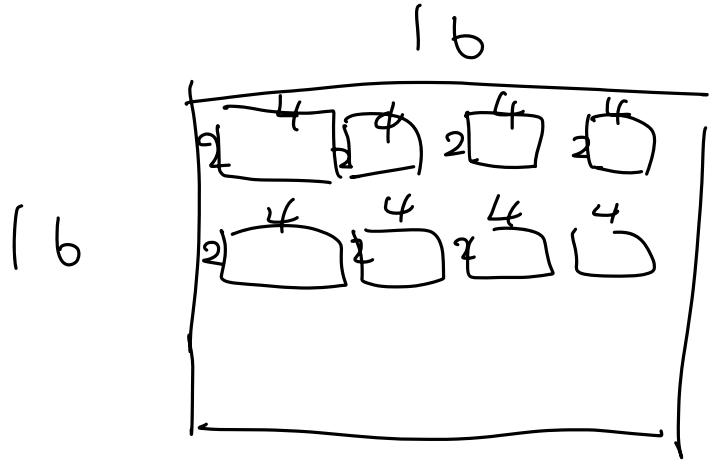

② How many TPU Tensor Cores?

16 x 16 x 1 = 256

TPU v5e has 1 TPU core per chip

③ Total FLOPs/sec for the whole pod

$1.97 \times 10^{14} \times 16 \times 16 = 504.32 \times 10^{14} = 5.04 \times 10^{16}$

④ Total HBM = 16 GB x 16 x 16

= 4096 GB = 4 TB

Question 3: PCIe operational intensity

$A = \text{bfloat}[D, F]$

$X \text{ Activations} = \text{bfloat}[B, D]$

on DRAMs

$$\text{FLOPs} = BF(\underbrace{D}_{\text{mul}} + \underbrace{D-1}_{\text{add}})$$

$$= BF(2D-1) = 2BDF - BF$$

$$\simeq 2BDF \checkmark$$

$$\text{FLOPs} = 2BDF \overset{F=4D}{=} 2BD \cdot 4D = 8BD^2$$

Load write

$$\text{Bytes} = 2(\underbrace{DF + BD}_{\text{Load}}) + \underbrace{2BF}_{\text{write}}$$

$$\overset{F=4D}{=} 2(4D^2 + BD) = 8D^2 + 2BD \simeq 8D^2$$

$$\frac{\text{FLOPs}}{\text{Bytes}} = \frac{8BD^2}{8D^2} = B$$

$$AI_{hardware} = \frac{\text{Assume single chip } 4 \times 9.2 \times 10^{14}}{1.5 \times 10^{6}}$$

$$= \frac{8 \times 9.2 \times 10^{4}}{1.5} = 49 \times 10^{4} \quad 61250$$

$$= 4.9 \times 10^{5}$$

Question 4: general matmul latency

Multiply int8[6384, 4096] by int8[B, 4096]

Assumes: 1 TPUv5e (chip)

1. ✓

Time = max(bytes load/write, FLOPs time)

$$= \max\left(\frac{6384 \times 4096 + B \cdot 4096 + B \cdot 16384}{8.1 \times 10^{11}}, \frac{B \times 16384 \times 2 \times 4096}{3.94 \times 10^{14}} \right)$$

$$= \max\left(\frac{6.7 \times 10^{7} + 20480B}{8.1 \times 10^{11}} + \frac{1.3 \times 10^{8} B}{3.94 \times 10^{14}} \right)$$

Compute-bound when $\frac{1.3 \times 10^8 B}{3.94 \times 10^{14}} > \frac{6.7 \times 10^7 + 2 \times 10^4 B}{8.1 \times 10^{11}}$

$\Leftrightarrow \frac{1.3B}{3.94 \times 10^6} > \frac{6.7 \times 10^3 + 2B}{8.1 \times 10^7}$

$\Leftrightarrow 1.3 \times 8.1 \times 10^7 B > 6.7 \times 10^3 \times 3.94 \times 10^6 + 2 \times 3.94 \times 10^6 B$

$\Leftrightarrow 10.53 \times 10^7 B - 7.88 \times 10^6 B > 26.398 \times 10^9$

$\Leftrightarrow 105.3 \times 10^6 B - 7.88 \times 10^6 B > 2.6 \times 10^{10}$

$\Leftrightarrow 97.42 \times 10^6 B > 2.6 \times 10^{10}$

$\Leftrightarrow B > \frac{2.6 \times 10^4}{97.42}$

$\Leftrightarrow B > 268$

2. Assumes: bandwith $= 22 \times 8.1 \times 10^{11}$

Compute-bound when:

$\frac{1.3 \times 10^8 B}{3.94 \times 10^{14}} > \frac{6.7 \times 10^7 + 2 \times 10^4 B}{22 \times 8.1 \times 10^{11}}$

$\Leftrightarrow \frac{1.3 \times B}{3.94 \times 10^6} > \frac{6.7 \times 10^3 + 2B}{22 \times 8.1 \times 10^7}$

$$\Leftrightarrow \quad 1.3 \times 22 \times 8.1 \times 10^7 B > 3.94 \times 6.7 \times 10^9$$

$$+ 2 \times 3.94 \times 10^6 B$$

$$\Leftrightarrow \quad 231.66 \times 10^7 B > 26.398 \times 10^9$$

$$+ 7.88 \times 10^6 B$$

$$\Leftrightarrow \quad 2316.6 \times 10^6 B - 7.88 \times 10^6 B > 2.6398 \times 10^{10}$$

$$\Leftrightarrow \quad 2308.72 \times 10^6 B > 2.6398 \times 10^{10}$$

$$\Leftrightarrow \quad 2.30872 \times 10^9 B > 2.6398 \times 10^{10}$$

$$\Leftrightarrow \quad B > \frac{2.6398 \times 10^{10}}{2.30872 \times 10^9}$$

$$= \underline{11.43} \ \#$$

Question 5: ICI bandwidth

TPUv5e : 4x4 slice

bfloat[8, 128, 8192]

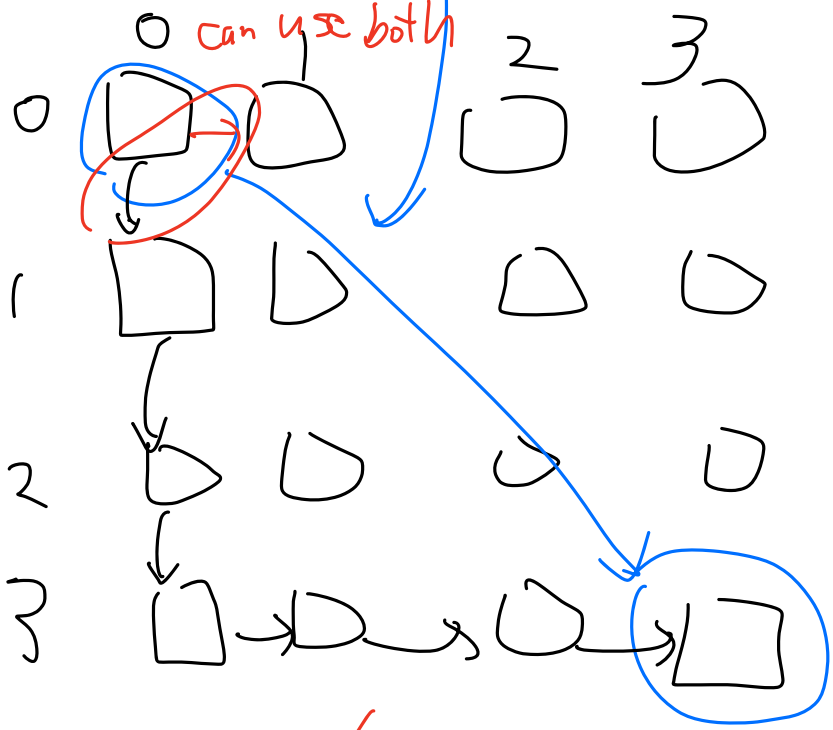

1. $6 \times 1\mu s$ = first byte latency
(hops)

2. Total time $= 6\mu s + \frac{2 \times 8 \times 128 \times 8192}{2 \times 4.5 \times 10^{10}}$

$= 6\mu s + \frac{16,777,216}{2 \times 4.5 \times 10^{10}}$

$= 6\mu s + \frac{1.6 \times 10^{7}}{2 \times 4.5 \times 10^{10}}$

$= 6\mu s + \frac{1.6}{2 \times 4.5 \times 10^{3}}$

$= 6 \times 10^{-6}$ sec $+ \frac{0.35 \times 10^{-3}}{2}$ sec

$= 6 \times 10^{-6}$ sec $+ \frac{3.5 \times 10^{-4}}{2}$ sec

$= 356 \mu s$

181

Question 6:

Matrix A: int8[128x1024, 128x1024]

Sharded evenly across TPUv5e 4x4 slice but they are on host's DRAM on each chip.

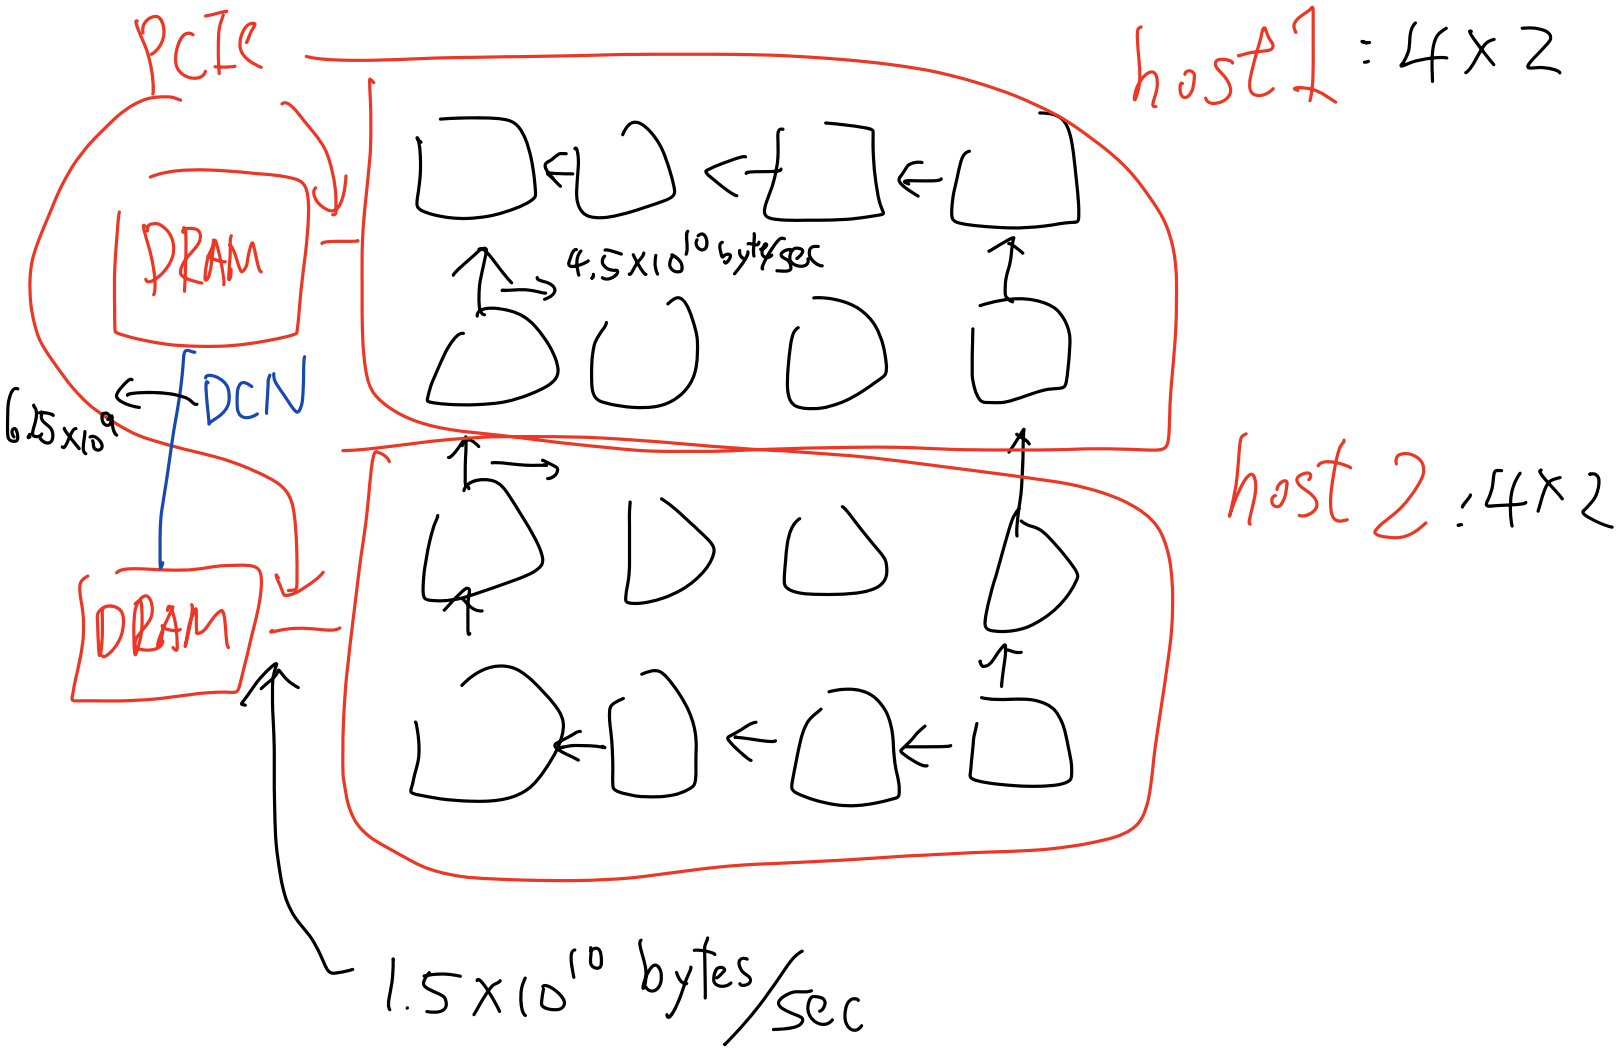

First: We can choose either:

① Copy half of the data through DCN to host 1's DRAM and then copy things to HBM from DRAM

or ② Each host copy its data to TPU chips and then utilize ICI to copy data

Option ①: DCN throughput is slow, so use ②

Steps: (will take the max over those numbers)

1. Copy data from DRAM to HBM
2. Transfer data from all chips to TPU [0, 0] 2^7 2^{10}
3. Compute int8[2^{11} x 2^{10}, 128x1024] bfloat16[8, 128x1024]

1. $128 \times 1024 \times 128 \times 1024 = 2^{34} = 2^4$ GB
$= 16$ GiB

$\frac{16\text{ GiB}}{2} = 8\text{ GiB} =$ data on each host

Time to load from DRAM to HBM (to all TPU chips) through PCIe

$= \frac{8\text{ GiB}}{16 \times 1.5 \times 10^{10}\text{ bytes/sec}} = 0.035\text{ sec} = 35\text{ ms}$

2. Time to send data from HBM to TPU{0,0} through ICI;

15 GiB because TPU{0,0} needs to receive 15 GiB

$\frac{15\text{ GiB}}{2 \times 4.5 \times 10^{10}} = \frac{15\text{ GiB}}{9 \times 10^{10}\text{ bytes/sec}}$

$= 0.178\text{ sec}$
$= 178\text{ ms}$

3. Time to load and write to MXU

Load int8[128x1024, 128x1024] = $\frac{128 \times 1024 \times 128 \times 1024}{8.1 \times 10^{11}} = 0.021$ $= 21$ ms

Load bf16[8, 128x1024] = I think the problem assume it is already in TPU HBM

Write bf16[8, 128x1024] $= 2 \times (8 \times 128 \times 1024)$
$= 2 \times 2^3 \times 2^7 \times 10^{10}$

$= 2\,MiB$

$$\frac{2\,MiB}{8.1 \times 10^{11}\ (\text{HBM speed})} = \text{negligible}$$

3. Compute time:

$$\frac{2BDF}{1.97 \times 10^{14}\ \text{FLOP/sec}} = \frac{2 \times (8 \times 128 \times 1024 \times 128 \times 1024)}{1.97 \times 10^{14}\ \text{FLOPs/sec}}$$

$$= \frac{2 \times (2^3 \times 2^7 \times 2^{10} \times 2^7 \times 2^{10})}{1.97 \times 10^{14}}$$

$$= \frac{2^{38}}{1.97 \times 10^{14}}$$

$$= \frac{256\ GiB}{1.97 \times 10^{14}} = 0.001395$$

$$= 1.3\,ms$$

Total = max stage

$= 178\,ms$